

Quiz(20%)

10.19 Thursday

13:00pm--14:00pm

No workshop this week

Unit 8:

Regions of Memory: Global, and Local Space

```

Offset(h) 00 01 02 03 04 05 06 07 08 09 0A 0B 0C 0D 0E 0F
00000660 48 65 6C 6C 6F 2C 20 77 6F 72 6C 64 2E 36 34 30 Hello, world.640
00000670 30 39 30 39 31 36 35 30 30 30 31 39 36 32 33 0D 090916500019623.
00000680 0A 3A 31 30 31 45 35 30 30 30 39 30 39 33 36 35 .:101E5000909365
00000690 30 30 38 30 39 33 36 34 30 30 32 46 35 46 33 46 00809364002F5F3F
000006A0 34 46 38 30 39 31 36 36 30 31 45 46 0D 0A 3A 31 4F80916601EF...1
000006B0 54 68 69 73 20 69 73 20 61 20 68 65 78 61 64 65 This is a hexade
000006C0 63 69 6D 61 6C 20 74 75 74 6F 72 69 61 6C 21 46 cimal tutorial!F
000006D0 38 39 34 45 31 39 39 33 36 0D 0A 3A 31 30 31 45 894E19936...101E
000006E0 37 30 30 30 00 01 02 03 04 05 06 07 08 09 0A 0B 7000.....
000006F0 0C 0D 0E 0F 10 11 12 13 14 15 16 17 18 19 1A 1B .....
00000700 1C 1D 1E 1F 20 21 22 23 24 25 26 27 28 29 2A 2B .... !"#%&'()*+
00000710 2C 2D 2E 2F 30 31 32 33 34 35 36 37 38 39 3A 3B ,-. /0123456789;;
00000720 3C 3D 3E 3F 40 41 42 43 44 45 46 47 48 49 4A 4B <=>?@ABCDEFGHJK
00000730 4C 4D 4E 4F 50 51 52 53 54 55 56 57 58 59 5A 5B LMNOPQRSTUVWXYZ[
00000740 5C 5D 5E 5F 60 61 62 63 64 65 66 67 68 69 6A 6B \]^_`abcdefghijklmnopqrstuvwxyz
00000750 6C 6D 6E 6F 70 71 72 73 74 75 76 77 78 79 7A 7B lmnopqrstuvwxyz{
00000760 7C 7D 7E 7F 80 81 82 83 84 85 86 87 88 89 8A 8B |}~.€.,f,,t$^&S(
00000770 8C 8D 8E 8F 90 91 92 93 94 95 96 97 98 99 9A 9B &Z...''''*._~m$>
00000780 9C 9D 9E 9F A0 A1 A2 A3 A4 A5 A6 A7 A8 A9 AA AB æ.žŸ ¡ÇE=¥!$"©ª«
00000790 AC AD AE AF B0 B1 B2 B3 B4 B5 B6 B7 B8 B9 BA BB ¬.®¯°±²³´µ¶·¸¹º»
000007A0 BC BD BE BF C0 C1 C2 C3 C4 C5 C6 C7 C8 C9 CA CB %&#2:ÁÀÃÄÅÆÇÈÉÊË
000007B0 CC CD CE CF D0 D1 D2 D3 D4 D5 D6 D7 D8 D9 DA DB ÌÍÎÏÐÑÒÓÔÕ×ØÙÚ
000007C0 DC DD DE DF E0 E1 E2 E3 E4 E5 E6 E7 E8 E9 EA EB ÛÜÝÞßàáâãäåæçèéêë
000007D0 EC ED EE EF F0 F1 F2 F3 F4 F5 F6 F7 F8 F9 FA FB ìíîïðñòóôõö÷øùúû
000007E0 FC FD FE FF B3 39 43 0D 0A 3A 31 30 31 45 44 30 uyby39C...101ED0
000007F0 30 30 35 37 30 30 45 38 39 35 33 32 39 36 30 32 005700E895329602

```

Hexadecimal

Binary / Decimal

- Assume a memory cell consists of 8 bits (one byte).
- We could always give the value of those 8 bits when discussing memory contents. (painful)
- We could convert them into decimal. (also painful)
 - E.g. $11111001_b = 249_d$

Hexadecimal : numbers to the base 16

binary \rightarrow 1 1 1 1 1 0 0 1
 decimal \rightarrow 15 9
 Hex \rightarrow F 9

Two Hex digits represent the value contained in a single byte (8 bits).

decimal	hexadecimal
0	0
1	1
2	2
3	3
4	4
5	5
6	6
7	7
8	8
9	9
10	A
11	B
12	C
13	D
14	E
15	F

- Consider an unsigned 16-bit integer.
- In Decimal, the conversion is even worse than for 1 byte.

2^{15}	2^{14}	2^{13}	2^{12}	2^{11}	2^{10}	2^9	2^8	2^7	2^6	2^5	2^4	2^3	2^2	2^1	2^0
1	1	0	1	0	1	1	1	1	0	1	1	1	0	0	1

$$32768 + 16384 + 4096 + 1024 + 512 + 256 + 128 + 32 + 16 + 8 + 1 \\ = 55225$$

Hex ⁶

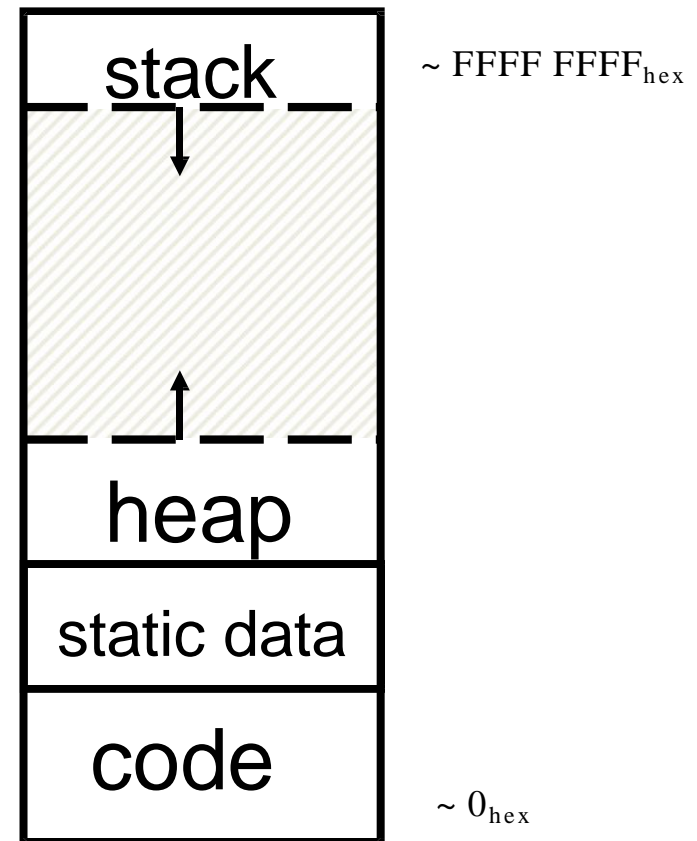
- With hexadecimal it is MUCH easier.

8	4	2	1	8	4	2	1	8	4	2	1	8	4	2	1
1	1	0	1	0	1	1	1	1	0	1	1	1	0	0	1
D				7				B				9			

Memory Space

Typical C Memory Management

- **stack**
 - local variables, parameters, return address
 - grows downward
- **heap**
 - space requested via e.g. `malloc()` ;
 - resizes dynamically, grows upward
 - data lives until deallocated by programmer
- **static data**
 - variables declared outside any functions, does not grow or shrink i.e. globals
- **code**
 - loaded when program starts, normally does not change



STATIC DATA REGION

Local Variables in 'C'¹⁰

```
int age = 0x7; // hex
char name [7] = "louise";
int salary = 0x7654; // hex
```

```
void doIncrement()
{
    int increment = 0x10; // hex
    salary = salary + increment;
}
```

```
void main()
{
    char reverse [7];
    doIncrement();
    printf("salary plus increment = %d\n", salary);
}
```

“increment” is local to the body of “doIncrement”. This means it can only be referred to within “doIncrement”.

If in the “main” procedure we referred to **“increment”**, this would result in a compile-time error.

Global Variables in 'C'11

```
int age = 0x7; // hex
char name [7] = "louise";
int salary = 0x7654; // hex


void doIncrement(){
    int increment = 0x10; // hex
    salary = salary + increment;
}

void main() {
    char reverse [7];
    doIncrement();
    printf("salary plus increment = %d\n", salary);
}
```

Quiz¹²

```
int  y = 0;
int  array[10];
int  *p;

int main() {
    int x;
    int *pt = NULL;
    p=malloc(sizeof(x));
    return 0;
}
```



~ FFFF FFFF_{hex}

~ 0_{hex}


Memory allocation for global ¹³ variables

- Assumptions:
 - The target machine has:
 - A 32-bit word size (i.e. the machine can load/store 4 bytes with one operation)
 - Memory is addressed on a per-byte basis

Some globals ¹⁴

- The globals

```
int age = 0x7; // hex
char name [7] = "louise";
int salary = 0x7654; // hex
```

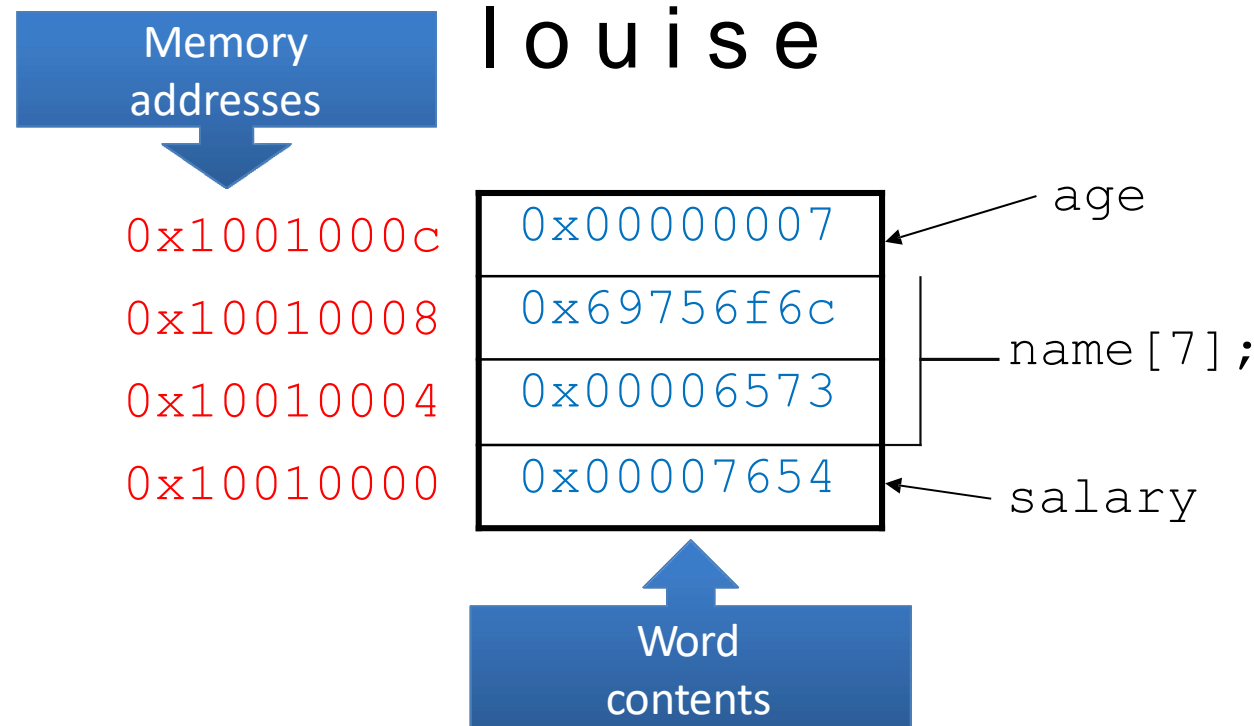


Our C program's global
declarations

- Will be laid out in memory as follows :

Remember:
4 bytes/word of memory
2 hex characters per byte

```
int age = 0x7;  
char name [7] = "louise";  
int salary = 0x7654;
```



'l'	'o'	'u'	'i'	's'	'e'	'\0'
6c	6f	75	69	73	65	0

ASCII codes

With byte addressing

Why `name[7]` takes up 8 bytes memory????

Hints:

The word-size of this machine is 4 bytes

0x1001000c	0x07
0x1001000b	0x00
0x1001000a	0x00
0x10010009	0x00
0x10010008	0x6c
0x10010007	0x6f
0x10010006	0x75
0x10010005	0x69
0x10010004	0x73
0x10010003	0x65
0x10010002	0x00
0x10010001	0x00
0x10010000	0x54
0x1000FFFF	0x76
0x1000FFFE	0x00
0x1000FFFD	0x00

l
o
u
i
s
e

LOCAL VARIABLES & STACK

How is space allocated for Local Variables?

- Local variables only exist at runtime when the function, procedure or method containing them is being called.
- We need some mechanism that allows us to
 - allocate space when a function, procedure or method is called, and
 - deallocate space (so it can be reused) when we return from a procedure

Stacks

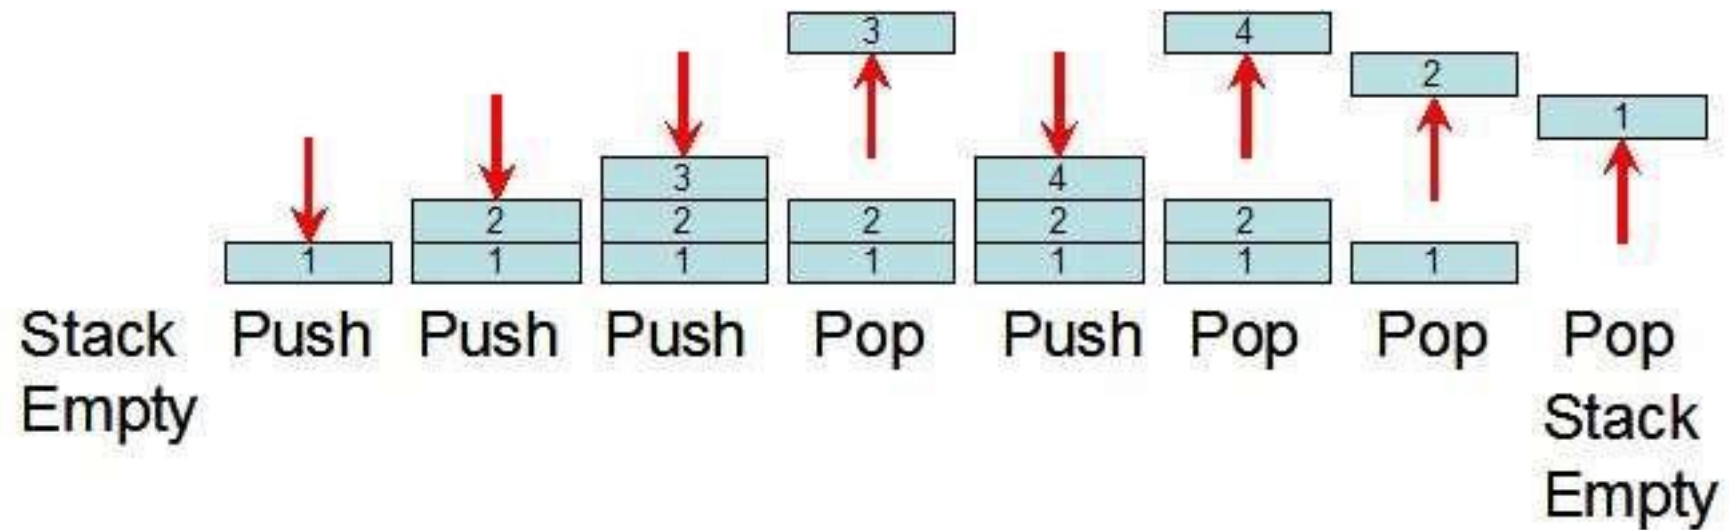
- What is a Stack?
 - A stack is a data structure of ordered items such that items can be inserted and removed only at one end.



Stacks (cont)

- What can we do with a stack?
 - `push` - place an item on the stack
 - `pop` - Look at the item on top of the stack and remove it
 - `isEmpty` - Reports whether the stack is empty or not

Popping and pushing



Quiz²²

- We can use a stack to reverse the letters of a string.
 - E.g. “Hello World!” -> “!dlroW olleH”
 - How?

Stacks (cont)

- Problem:
 - What happens if we try to pop an item off the stack when the stack is empty?
 - This is called a stack underflow.
 - The pop method needs some way of telling us that this has happened.

A SIMPLE INT STACK IN 'C'



pushing and popping (1)²⁵

push(x);	
push(y);	
x = pop();	
y = pop();	

sp	100
x	5
y	6

stack

99	??
98	??
97	??
96	??
...	??
0	??

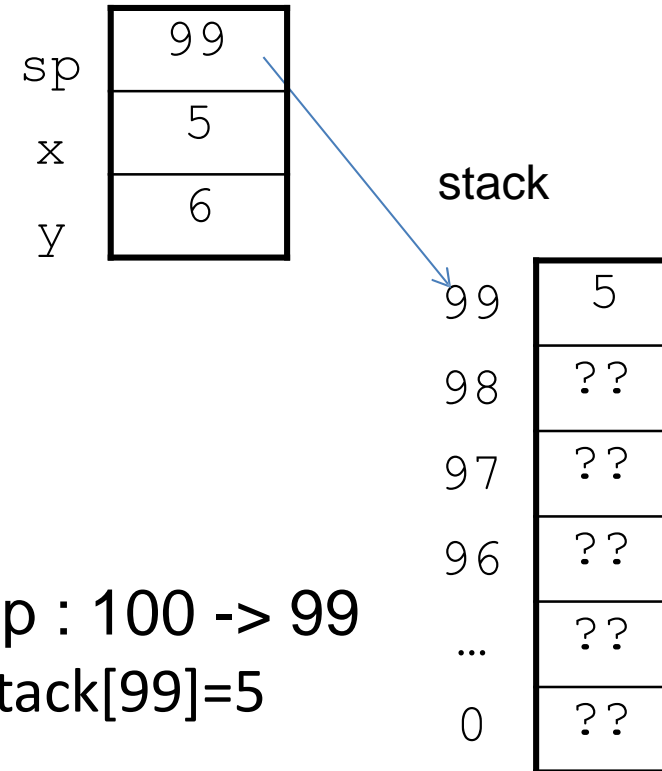
```
int stack[100];
int sp = 100; //Stack pointer: NOT a
               // pointer data type

void push (int value)
{
    sp = sp-1;
    if (sp < 0) return; //overflow
    stack[sp] = value;
}
```



pushing and popping (2)²⁶

push(x);	*
push(y);	
x = pop();	
y = pop();	



```
void push (int value)
{
    sp = sp-1;
    if (sp < 0) return; //overflow
    stack[sp] = value;
}
```

sp : 100 -> 99
stack[99]=5



pushing and popping (3)²⁷

push(x);	
push(y);	*
x = pop();	
y = pop();	

sp	99
x	5
y	6

stack

99	5
98	6
97	??
96	??
...	??
0	??

sp : 99 -> 98
stack[98]=6

```
void push (int value)
{
    sp = sp-1;
    if (sp < 0) return; //overflow
    stack[sp] = value;
}
```



pushing and popping (4)²⁸

<code>push(x);</code>	
<code>push(y);</code>	
<code>x = pop();</code>	*
<code>y = pop();</code>	

sp	98
x	5
y	6

stack

99	5
98	6
97	??
96	??
...	??
0	??

```
int pop()
{
    if (sp >= 100) return -1; //stack underflow!
    int value = stack[sp];
    sp = sp + 1;
    return value;
}
```

value = stack [98] = 6

sp: 98 -> 99



pushing and popping (5)²⁹

push(x);	
push(y);	
x = pop();	
y = pop();	*

sp	99
x	5
y	6

stack

99	5
98	??
97	??
96	??
...	??
0	??


```
int pop()
{
    if (sp >= 100) return -1; //stack underflow!
    int value = stack[sp];
    sp = sp + 1;
    return value;
}
```

value = stack [99] = 5

sp: 99 -> 100

The run-time Stack

- Assume the stack begins at the high-end of the data portion of memory.
- As items are added, the stack pointer is decremented.
 - “grow downwards” towards the program section of memory.




How Local Variables are pushed³¹ onto the stack?

```
int age = 0x7; // hex
char name [6] = "louise";
int salary = 0x7654; // hex
```

```
void doIncrement()
{
    int increment = 0x10; // hex
    salary = salary + increment;
}
```

```
void main()
{
    char reverse [7];
    doIncrement();
    printf("salary plus increment = %d\n", salary);
}
```



Making room for `reverse[7]`³²


- Assume the stack pointer is set to 1000
- There are 7 bytes in `reverse`, but we will round this up to a word boundary
 - hence we must reserve 8 bytes on the stack ...
 - `sp = 992` after allocating space for `reverse[7]`

1000	??
999	??
998	??
997	??
996	??
995	??
994	??
993	??
992	??
991	??

Calling doincrement

```
int age = 0x7; // hex
char name [6] = "louise";
int salary = 0x7654; // hex
```

```
void doIncrement()
{
    int increment = 0x10; // hex
    salary = salary + increment;
}
```

```
void main()
{
    char reverse [7];
    doIncrement(); 
    printf("salary plus increment = %d\n", salary);
}
```

Calling doincrement (2)³⁴

```
int age = 0x7; // hex
char name [6] = "louise";
int salary = 0x7654; // hex
```

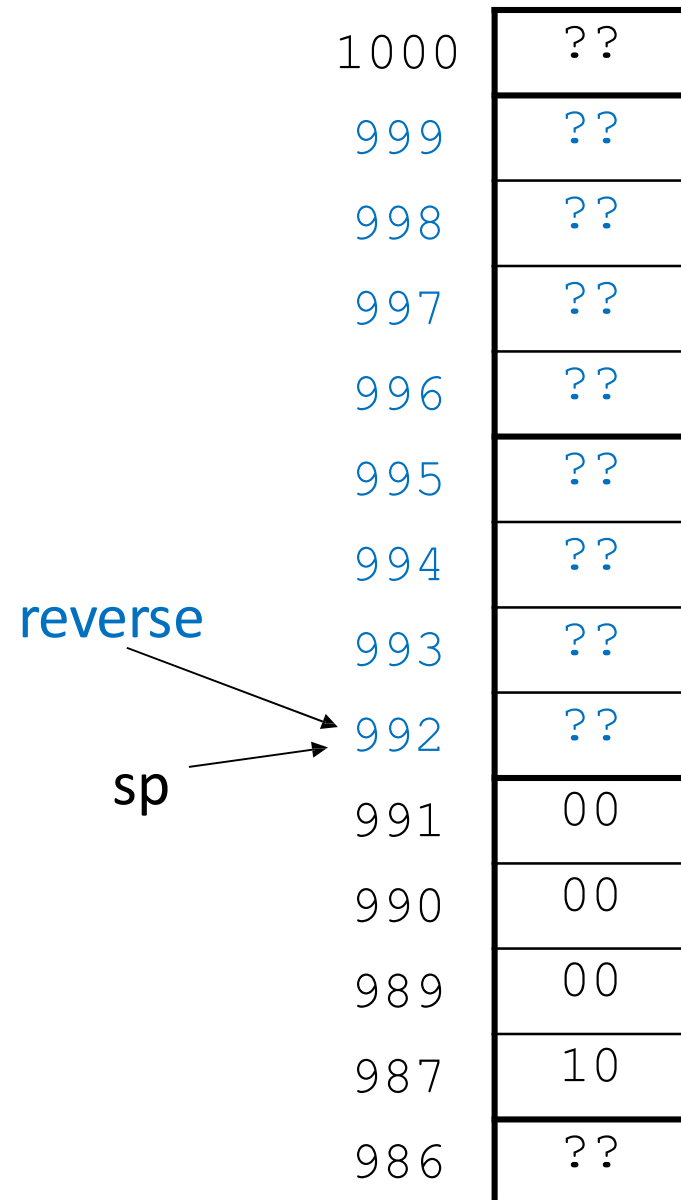
```
void doIncrement()
{
    int increment = 0x10; // hex
    salary = salary + increment;
}
```



```
void main()
{
    char reverse [7];
    doIncrement();
    printf("salary plus increment = %d\n", salary);
}
```

Calling doIncrement³⁵

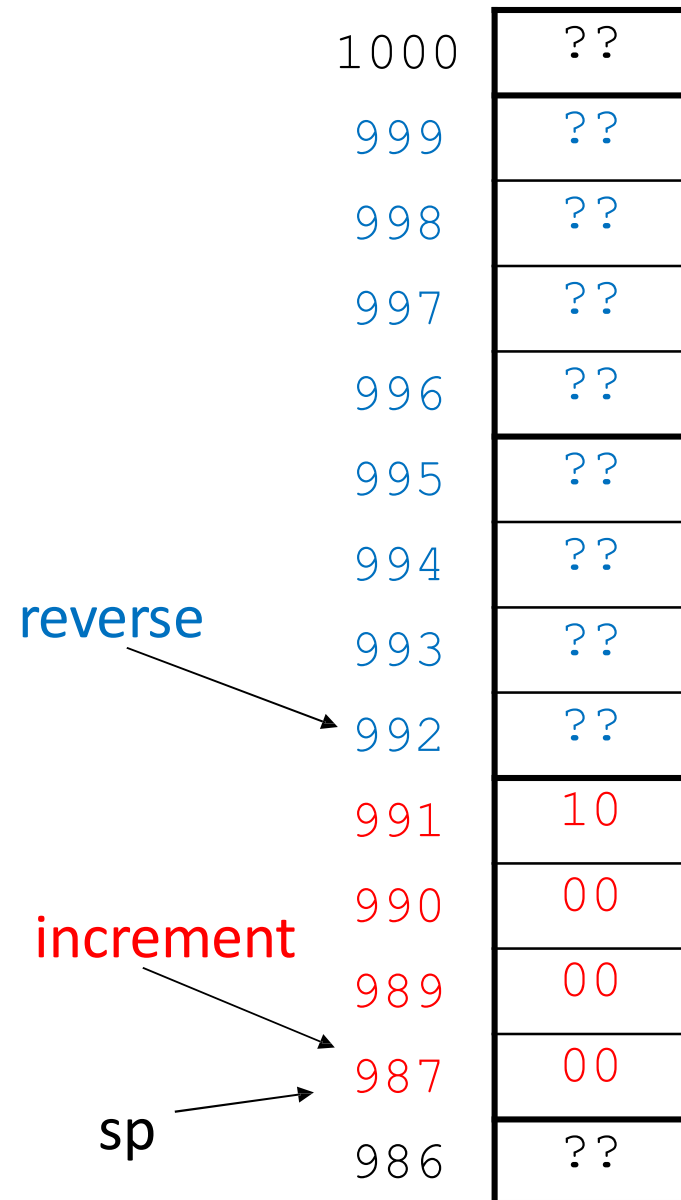
- `doIncrement` contains a local variable,
`int increment = 0x10`
- We need to allocate space (4 bytes) for this on the stack.



Calling doIncrement³⁶

- `doIncrement` contains a local variable,
`int increment = 0x10`
 - We need to allocate space (4 bytes) for this on the stack.

If `doIncrement` had any parameters, we would also have to allocate space for them on the stack.



Managing procedure calls³⁷

- Whenever a subroutine (procedure, function) is called, some storage must be set aside for
 - local variables
 - parameters (or arguments)
 - management information i.e. the return address

```
void main()  
{  
    char reverse [7];  
    foo(20);  
    printf("salary plus increment = %d\n", salary);  
}
```

the call

At the end of the call, control returns to here

The diagram illustrates the control flow of a function call. A box highlights the call to `foo(20);` within the `main` function. An arrow points from this call to the text "the call". Another arrow points from the `printf` statement back to the `foo(20);` call, with the text "At the end of the call, control returns to here" below it.

Stack frame

- A stack frame is an area that stores:

Copies of values of parameters
All local variables
Return result (if any)
Return address

- When a function is called, a new stack frame is created and pushed on the stack.
- And every time we exit a function, the stack frame is popped off the stack.

Creating Stack Frames


```
void d (int p)
{
}

void c (int o)
{ d(o);
}

void b (int n)
{ c(n);
}

void a (int m)
{ int a = m+1; b(a);
}

main ()
{
    a(0);
    printf("hello");
}
```



Creating Stack Frames


```
void d (int p)
{
}
```

```
void c (int o)
{ d(o);
}
```

```
void b (int n)
{ c(n);
}
```

```
void a (int m)
{ int a = m+1; b(a);
}
```

```
main ()
{
    a(0);
    printf("hello");
}
```



The return address in the stack frame of `a(0)` held the Address of the next instruction in main. It is the address of the instruction from which execution should be resumed once `a(0)` has finished executing.



Unit 9:


Dynamic storage structures and memory management

Aims of this unit

- At the end of this unit you will understand :
 - The heap :
 - the area of memory used to dynamically allocate (e.g. `malloc`) and free memory cells, as used in dynamic data structures such as linked lists
 - How the heap space can be managed, both by
 - Automatic garbage collection (used by the language system)
 - System calls (used directly by the programmer)

Memory Management Revisit ³

- **stack**
 - Variable size, variable content (local variables etc.)
 - Used for managing function calls and returns
- **heap**
 - Dynamically allocated objects and data structures
 - data lives until deallocated
- **static data**
 - Fixed size, fixed content, allocated at compile time
- **code**
 - loaded when program starts, does not change



The Heap (A Typical ⁴ Architecture)

- Contains a linked list of used and free cells (blocks)
- New cells are created as needed, e.g., using `new` (in java) or `malloc` (in C)
- Redundant (no in-used) cells should be released
 - Question: How can redundant cells be released?

Memory blocks can be different ⁵ in sizes

- In a hybrid Java/'C' style :

```
struct IntCell
{
    int data;          //4 bytes
    struct IntCell* next;
    //32 bit addressing - 4 bytes
};
```

```
IntCell ic = new IntCell();
//new creates an instance of
IntCell
```

Memory blocks can be different ⁶ in sizes (2)

```
struct PersonCell
{
    char name [12];           // 12 bytes
    int age;                  // 4 bytes
    struct IntCell* next;     // 4 bytes
};
```

- A PersonCell instance has $12 + 4 + 4 = 20$ bytes.

```
PersonCell pc = new PersonCell();
```

20 bytes allocated from the heap.


Memory Fragmentation -- ⁷

Assumptions

- Assume a tiny heap with 56 bytes
 - just to keep our examples simple!
- We have six variables : PC1, PC2 and PC3 of type `PersonCell` (20 bytes each)
- and IC1, IC2 and IC3 of type `IntCell` (8 bytes each).

How fragmentation can arise (1)⁸

```
PersonCell PC1 = new PersonCell();
```



Free space : 56


Bytes required : 20

After allocation, free
space left : 36

How fragmentation can arise (2)⁹

```
IntCell IC1 = new IntCell();
```


Free space : 36
Bytes required : 8
After allocation,
free space left : 28



How fragmentation can arise (3)¹⁰

```
PersonCell PC2 = new PersonCell();
```


Free space : 28
Bytes required : 20
After allocation,
free space left : 8



How fragmentation can arise (4)¹¹

```
IntCell IC2 = new IntCell();
```

Free space : 8
Bytes required : 8
After allocation,
free space left : 0




How fragmentation can arise (5)¹²

```
free(PC1); //not possible in Java!
```

Free space : 0

Bytes released : 20

After deallocation,
free space left : 20




How fragmentation can arise (6)¹³

```
IntCell IC3 = new IntCell();
```

Free space : 20

Bytes required : 8

After allocation, free
space left : 12




How fragmentation can arise (7)¹⁴

`free(IC2);`

Free space : 12

Bytes released : 8

After deallocation,
free space left : 20





How fragmentation can arise (8)¹⁵

```
PersonCell PC3 = new PersonCell(); ???
```

Free space : 20

Bytes required : 20

But cannot allocate as
the 20 bytes are not
contiguous



'C' system calls : a reminder

- `malloc()`: Allocates raw, uninitialized memory from heap (malloc -> memory allocate)
- `sizeof : malloc` needs to know how many bytes to allocate;
 - given a C type as parameter, this returns how many bytes a variable of that type takes up.
- `free()` : deallocates memory

The Heap : Freeing Unwanted¹⁸ Cells

- Dangerous:
 - With `free` the programmer can free an in-use cell. (i.e. the cell is shared by two or more data structures)
- Resulting errors are
 - hard to detect
 - may be unnoticed
- These are some reasons why perhaps it is better for cells to be freed by the system rather than the programmer



GARBAGE COLLECTION

Garbage in Java

- Objects (data + methods) take up space
- Java automatically figures out which objects are not being used
- Cells (objects) in dynamic data structures in Java become “garbage” when no-one is pointing at them anymore.

Garbage in Java (2)

```
PersonCell pc = new PersonCell();  
    //Object 1  
PersonCell pc2 = pc;  
//pc2 refers to object 1 too  
pc.age = 20;  
store_to_a_file(pc)  
  
//now finished with the object 1  
  
pc = new PersonCell(); //object 2  
pc2 = new PersonCell(); //object 3  
  
// pc and pc2 now all point to a new  
object, so object 1 becomes garbage
```

The Heap : Garbage Collection

- Garbage Collector
 - Finds redundant cells automatically, so programmer need not worry
 - Called as needed, or in background
- Many different approaches
 - But there is a performance and space overhead
 - Used by (e.g.) Java, Python, Scala...

The Need for Garbage²³ Collection

- At a given moment, the heap contains:
 - cells which are live (in-use)
 - cells which are redundant (formerly in use) :
 - cells which are known to be free
- If creation of an object fails:
 - look for redundant cells
 - designate them to be free
 - try to create the object again

Mark & Sweep Garbage²⁴ Collection

- **Marking**
 - Mark objects found as in-use (live);
 - E.g. Simply count the number of references to an object
 - finally, free all unmarked objects.
- **Merging**
 - If no free cell is large enough, due to fragmentation of the heap, merge adjacent free cells.
- **Compaction**
 - If still no large enough cell, move live cells to bottom of heap, leaving large free area at top.
 - Compaction is slow, due to need to reevaluate pointers


Compaction in practice (1)²⁵

Free space : 20

Bytes required : 20

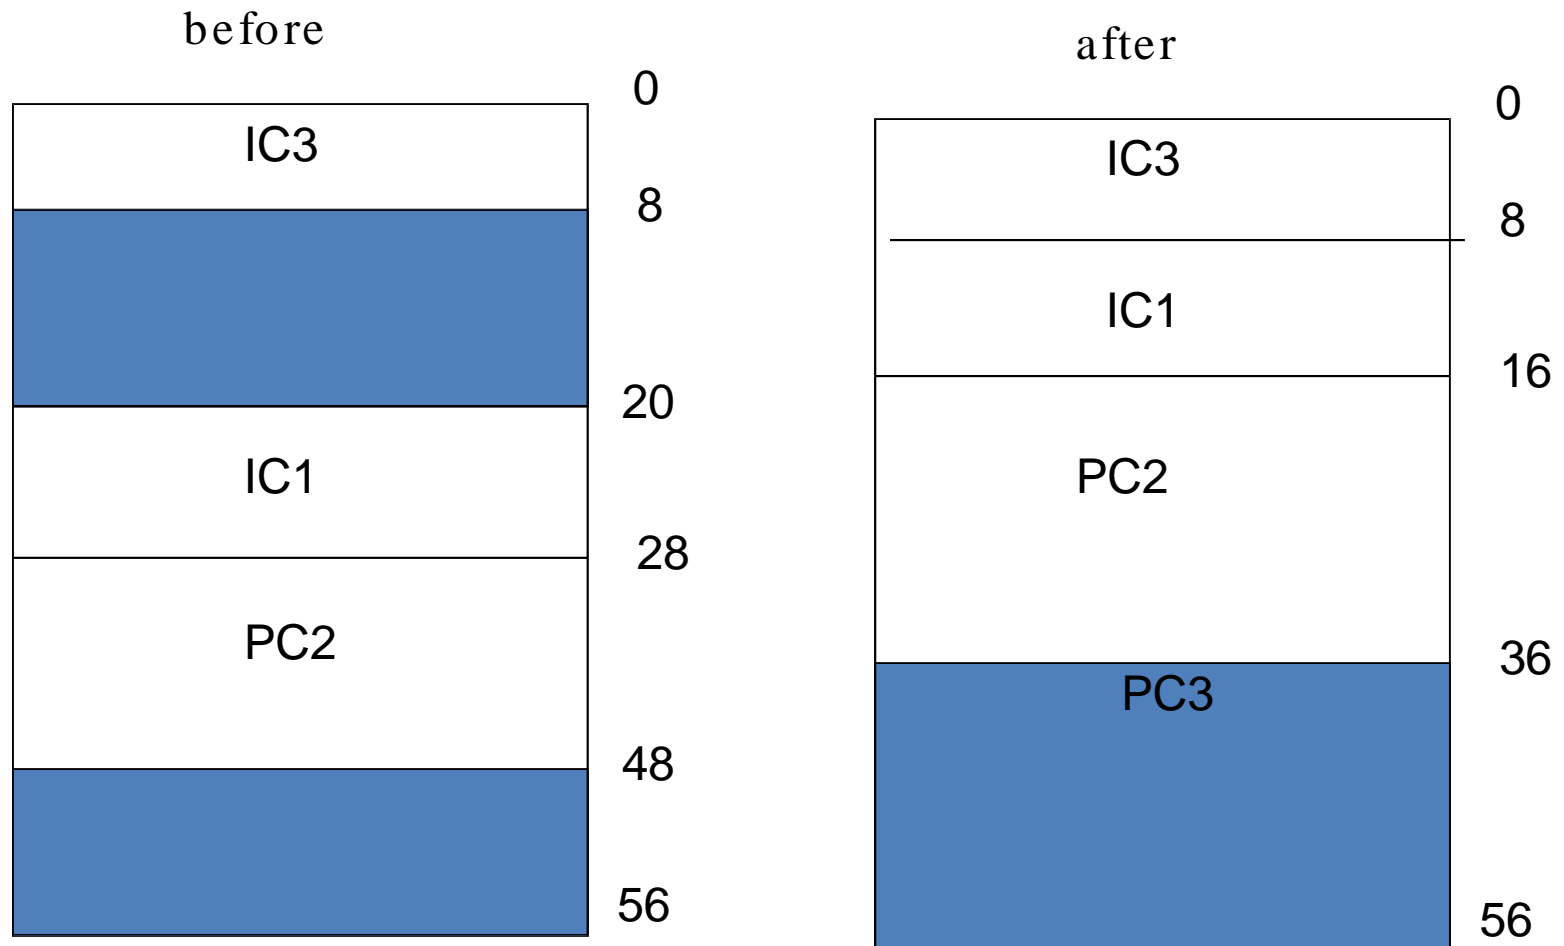
But cannot allocate as the
20 bytes are not
contiguous

IC3 points to HEAP[0];
IC1 points to HEAP[20];
PC2 points to HEAP[28];



```
PersonCell PC3 = new PersonCell(); ???
```

Compaction in practice (2)²⁶



```
PersonCell PC3 = new PersonCell();
```

Compaction in practice (3)²⁷

Before compaction :

IC3 points to HEAP[0];

IC1 points to HEAP[20];

PC2 points to HEAP[28];


These have to be updated to

IC3 points to HEAP[0];

IC1 points to HEAP[8];

PC2 points to HEAP[16];

```
PersonCell PC3 = new PersonCell();
```



Garbage Collection in Java²⁸

(cont)

- We can't explicitly call a “free” style method, but we can deliberately set a object to `null`.

```
PersonCell pc = new PersonCell(); //object 1
...
pc = null;           // pc no longer points at
                     //the space allocated for object 1
```

- Also, if the object is a local variable on the stack, once we have popped off the containing stack frame, the variable no longer exists and thus it is no longer pointing at the space.

'C' Malloc/Free Implementation

- In 'C', how might the system calls `malloc` and `free` manage the heap?
- Each cell of memory is preceded by a header that has two fields:
 - size of the cell and
 - a pointer to the next cell
- All free cells are kept in a linked list

Simple Implementation

- `malloc()` searches the free list for a cell that is big enough. If none is found, more memory is requested from the operating system.
- `free()` checks if the cells adjacent to the freed cell are also free
 - If so, adjacent free cells are merged (coalesced) into a single, larger free cell
 - Otherwise, the freed cell is just added to the free list

Choosing a cell in malloc()

- How do we choose which free cell to use?
 - **best-fit**: choose the smallest cell that is big enough for the request
 - **first-fit**: choose the first cell we see that is big enough
 - **next-fit**: like first-fit but remember where we finished searching and resume searching from there
- SCC students will learn more about this in SCC.211 Operating Systems

The END